



Learning Rules Across the Brain

Randall C. O'Reilly
University of Colorado Boulder
eCortex, Inc.

Ultimate Goal: Build a Brain that Builds Itself..


- Start with basic neural infrastructure...
- Add a lot of real world experience...
- Presto: instant human intelligence!
- How does human brain pull off this magic trick?



Kai (15mo)

Max (3.5yr)

Emer (?)



Learning Rules Across the Brain

Area	Learning Signal			Dynamics		
	Reward	Error	Self Org	Separator	Integrator	Attractor
<i>Primitive</i> Basal Ganglia	+++	---	---	++	-	---
Cerebellum	---	+++	---	+++	---	---
<i>Advanced</i> Hippocampus	+	+	+++	+++	---	+++
Neocortex	++	+++	++	---	+++	+++



+ = has to some extent ... +++ = defining characteristic – definitely has
 - = not likely to have ... --- = definitely does not have

Primitive, Basic Learning..



Area	Learning Signal				Dynamics	
	Reward	Error	Self Org	Separator	Integrator	Attractor
Primitive Basal Ganglia	+++	---	---	++	-	---
Cerebellum	---	+++	---	+++	---	---

- Reward & Error = most basic learning signals
(self organized learning is a luxury..)
- Simplest general solution to any learning problem is a *lookup table* = separator

Lookup Table & Pattern Separation



Cerebellar Error-driven Learning



Cerebellum =
Support Vector Machine

- Granule cells = high-dimensional encoding (separation)
- Purkinje/Olive = delta-rule error-driven learning
- Classic ideas from Marr (1969) & Albus (1971)


Cerebellum is Feed Forward

Feedforward circuit:

- Input (PN) -> granules -> Purkinje -> Output (DCN)
- Inhibitory interactions – no attractor dynamics
- Key idea: does delta-rule learning bridging small temporal gap:

$S(t-100) \rightarrow R(t)$

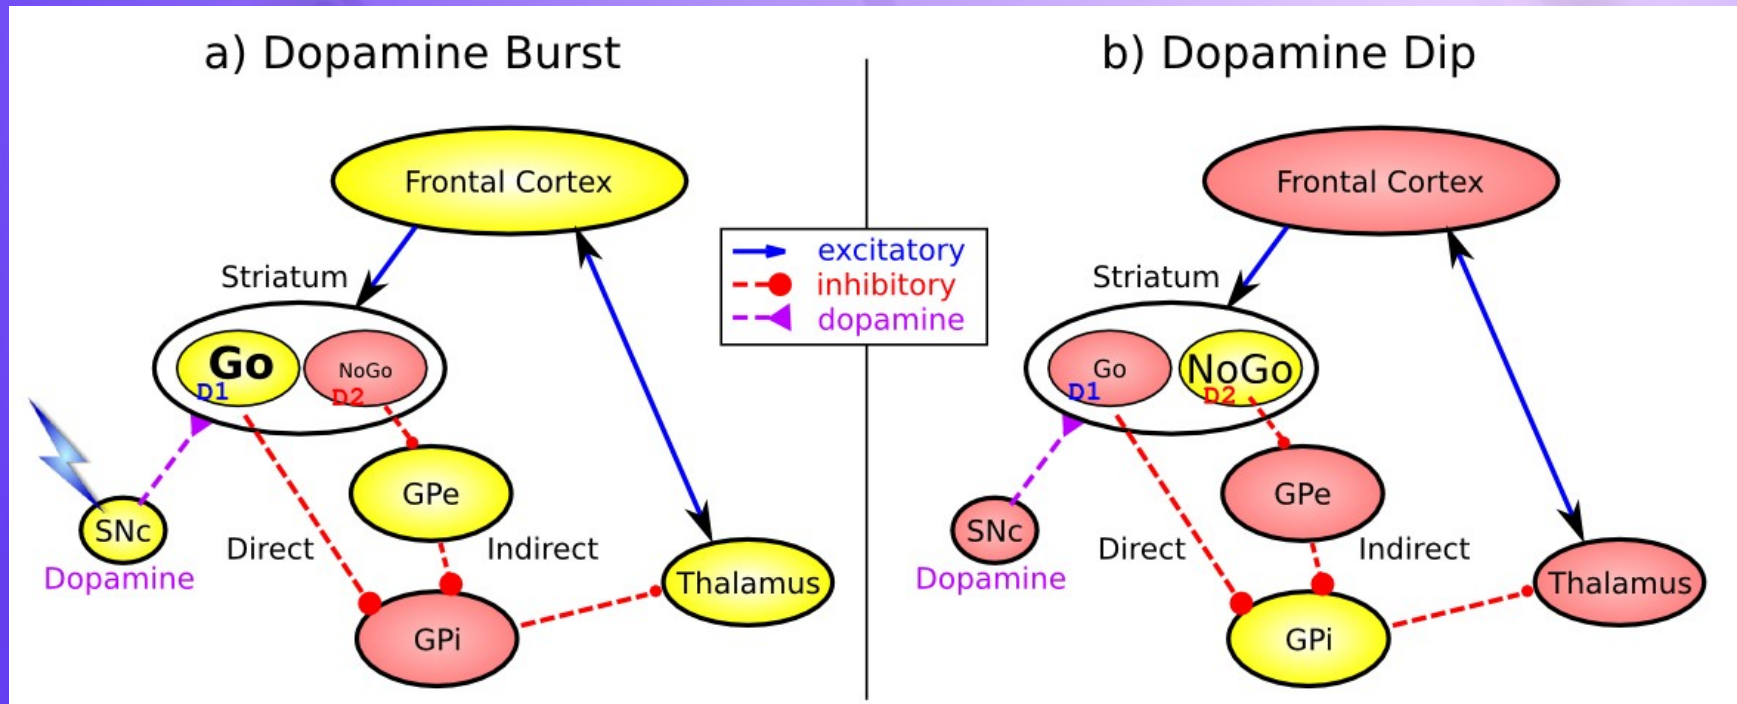
$\wedge \text{Error}(t+100)$



Bob Crimi

Basal Ganglia Reward Learning

(Frank, 2005...; O'Reilly & Frank 2006)



- Feedforward, modulatory (disinhibition) on cortex/motor (same as cerebellum)
- Co-opted for higher level cognitive control -> PFC


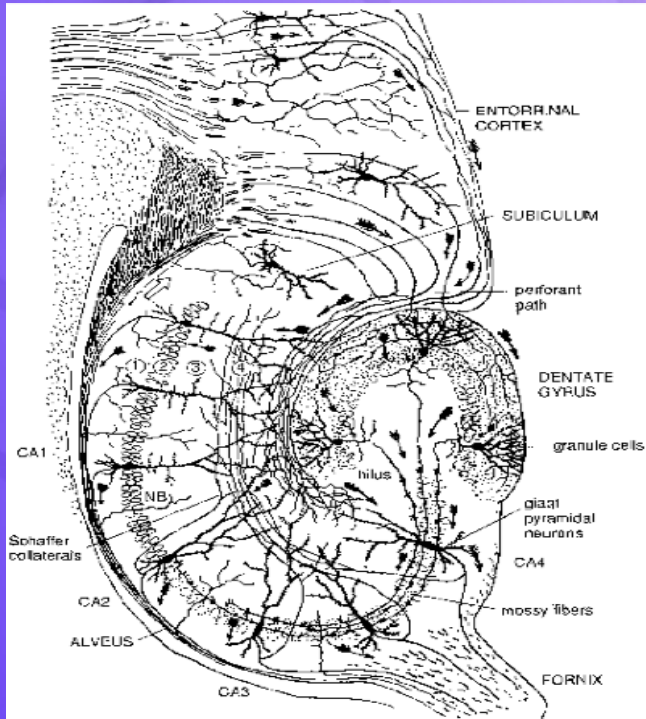
BG + Cerebellum Capacities

- Learn what satisfies basic needs, and what to avoid (BG reward learning)
 - And what information to maintain in working memory (PFC) to support successful behavior
- Learn basic Sensory -> Motor mappings accurately (Cerebellum error-driven learning)
 - Sensory -> Sensory mappings? (what is going to happen next..)

BG + Cerebellum Incapacities

- Generalize knowledge to novel situations
 - Lookup tables don't generalize well..
- Learn abstract semantics
 - Statistical regularities, higher-order categories, etc
- Encode episodic memories (specific events)
 - Useful for instance-based reasoning
- Plan, anticipate, simulate, etc..
 - Requires robust working memory

Hippocampal Episodic Memory



Area	Reward	Error	Self Org	Separator	Integrator	Attractor
Primitive Basal Ganglia	+++	---	---	++	-	---
Cerebellum	---	+++	---	+++	---	---
Advanced Hippocampus	+	+	+++	+++	---	+++


Hebbian Learning = Automatic

(O'Reilly & Rudy, 2001)

- Hippocampus is a constant “tape recorder” of your life
- You never know when something interesting might have happened 5 minutes ago..

Recall = Attractor

(content addressable memory..)



Recurrent connectivity is powerful but dangerous
– hippocampus is major source of epilepsy..

Neocortex!

Area	Learning Signal				Dynamics	
	Reward	Error	Self Org	Separator	Integrator	Attractor
<i>Primitive</i> Basal Ganglia	+++	---	---	++	-	---
Cerebellum	---	+++	---	+++	---	---
<i>Advanced</i> Hippocampus	+	+	+++	+++	---	+++
Neocortex	++	+++	++	---	+++	+++

- Integrates many different experiences to generate synthetic semantic knowledge (i.e., *common sense*), serves as the basis for generalization
- Integrative learning is necessarily slow!
(long life, neotony) (McClelland, McNaughton & O'Reilly, 1995)

Slow Integrative Learning

- Average: each term weighted by $1/N$
- Running average: $x(t) = x(t-1) + 1/\tau(s(t) - x(t))$
As $1/\tau$ gets smaller, window of integration gets larger
- Can't average over everything!!
 - Just gives you a big gamoosh blob!
- ***Integrative learning requires error-driven learning!***



Spike Based Error-Driven Learning

For 25 years:


- Biology (Hebbian) < - > Cognition* (Backprop)

* e.g., recognizing 100's of objects, learning to read..

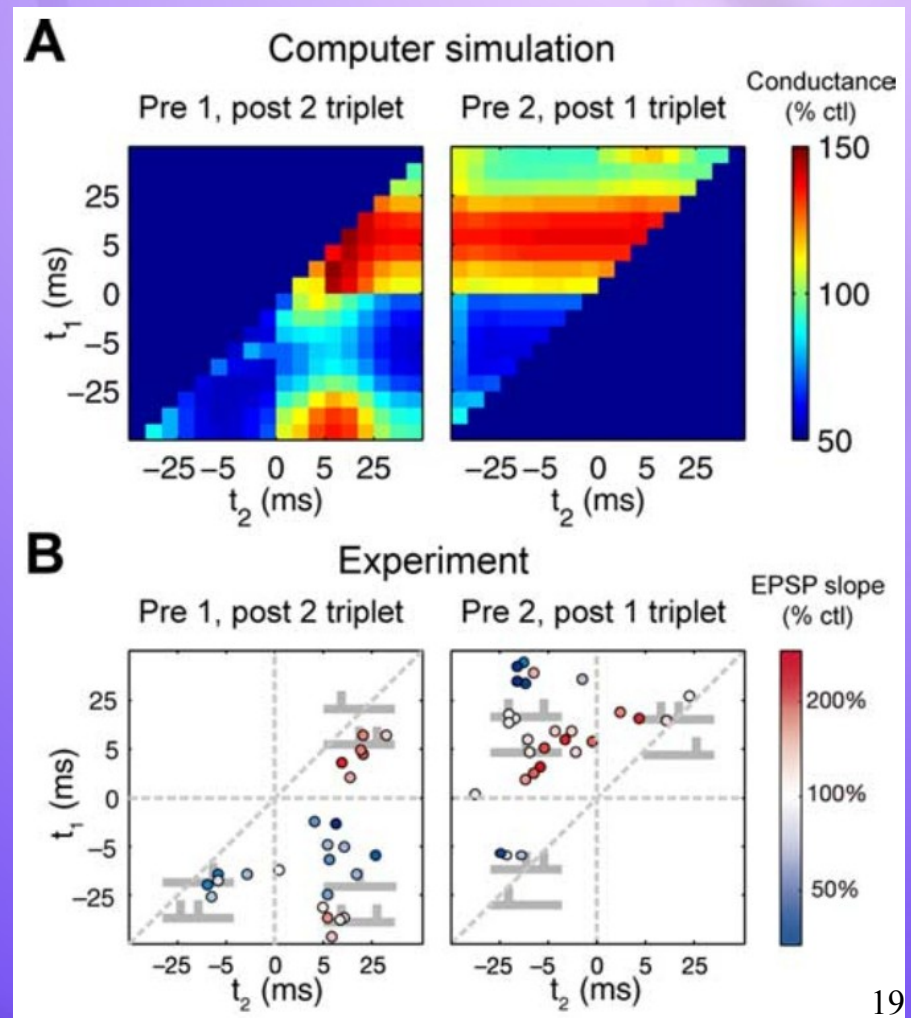
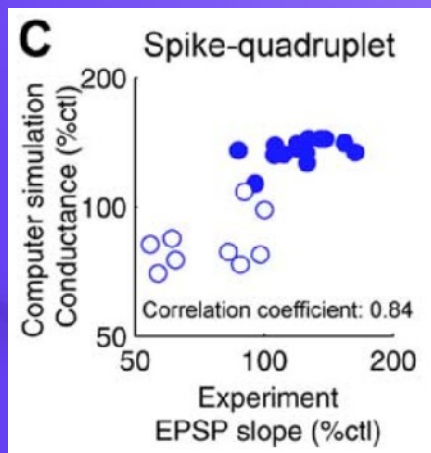

- Biology has advanced:
 - Spike Timing Dependent Plasticity (STDP)
 - Urakubo et al (2008) model: bottom-up brawn
- Cognition too:
 - Backprop -> Contrastive Hebbian Learning (CHL), Deterministic / Restricted Boltzmann Machine, GeneRec, Xie & Seung...
- **Can we unify them once and for all??**

Urakubo et al, 2008 Model


- Highly detailed combination of 3 existing strongly-validated models:



“Allosteric” NMDA Captures STDP (including higher-order and time integration effects)

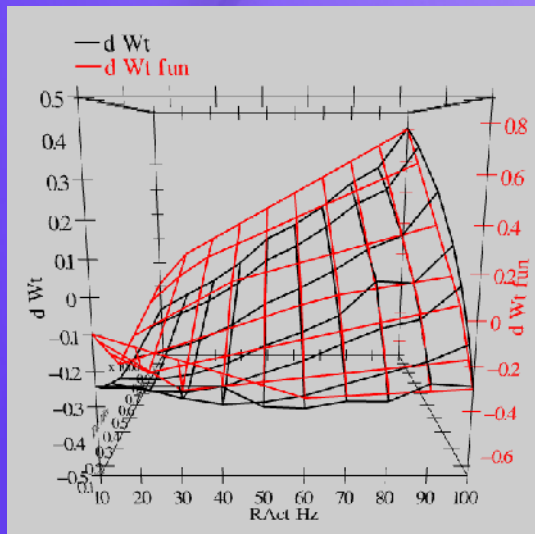


What About Real Spike Trains?




Extended Spike Trains = Emergent Simplicity


S = 100Hz



S = 50Hz




S = 20Hz




$$dW = f(\text{send} * \text{recv}) =$$

$$(\text{spike rate} * \text{duration})$$



$r = .894$


Linearized BCM



- Bienenstock, Cooper & Munro (1982) – BCM:
- **adaptive threshold Θ**
 - Lower when less active
 - Higher when more..

(homeostatic)

Fast Threshold Adaptation: Late Trains Early



Essence of Err-Driven: $dW = \text{outcome} - \text{expectation}$


Neocortical Bidirectional Connectivity (Attractors)



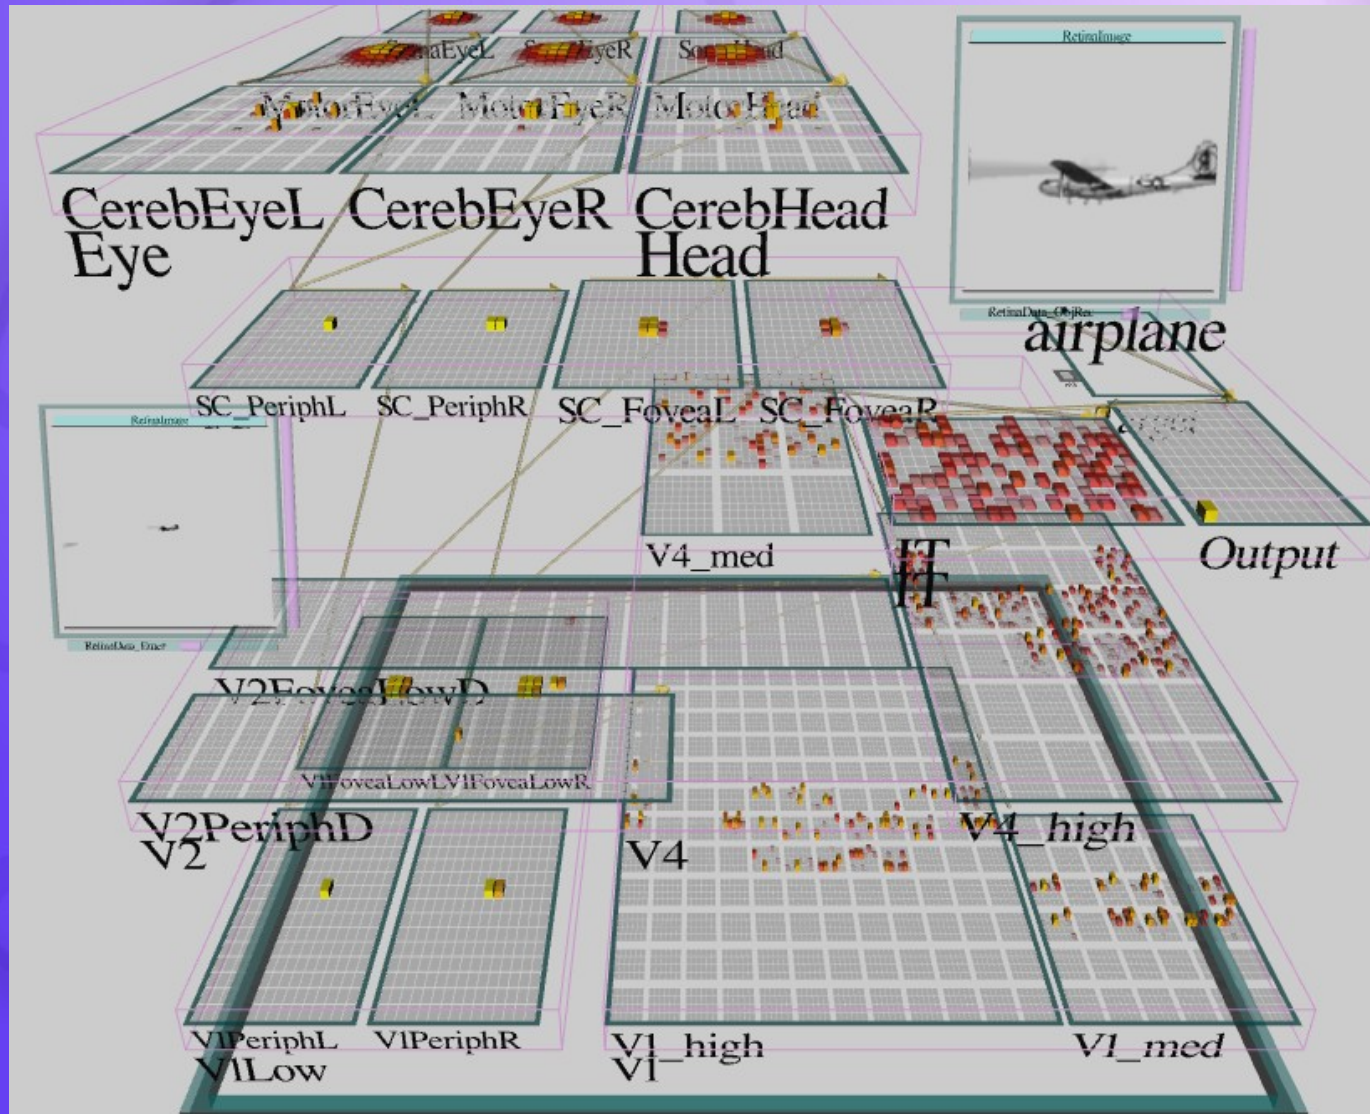
TAE CAT

- Constraint satisfaction: higher-level interpretation feeds back to constrain lower-level feature detectors
- Error-driven learning: conveys err signals

Emer Demo



Emer's Brain (Gaze + ObjRec)




Learning Rules Across the Brain

Area	Learning Signal				Dynamics	
	Reward	Error	Self Org	Separator	Integrator	Attractor
<i>Primitive</i> Basal Ganglia	+++	---	---	++	-	---
Cerebellum	---	+++	---	+++	---	---
<i>Advanced</i> Hippocampus	+	+	+++	+++	---	+++
Neocortex	++	+++	++	---	+++	+++

+ = has to some extent ... +++ = defining characteristic – definitely has
 - = not likely to have ... --- = definitely does not have

Leabra Tripartite Architecture



- Common neural mechs
- But specializations to resolve computational tradeoffs (e.g., fast, distinct learning vs. slow, integrative learning)
- Control flows from BG action selection to PFC updating to top-down bias

Thanks To

Rich Ivry – lots of fun discussion at Yali's!

CCN Lab

- Brad Aisa (Auditory)
- Tom Hazy (Learning)
- Seth Herd (Learning)
- Dave Jilk (Vision, eCortex)
- Kenneth Latimer (Vision)
- Brian Mingus (Motor)
- Wolfgang Pauli (Attention)
- Dean Wyatte (Vision)

Funding

- ONR – McKenna & Bello
- NIMH P50-MH079485
- NIMH R01-MH069597
- AFOSR
- DARPA

Extras


Object Recognition is Hard



- Large amount of shape variability within and between categories
- Huge amount of view-based variability (position, orientation, size, rotation)

Learning to Look is Non-Trivial

- 5 DOF -> 3 dimensions
- New gaze contingent on current positions
- Head & eye coordination



How Emer's Brain Works

1. Learning to Look

- How a “dumb” signal can be smart (Kawato et al)


1. Recognizing Objects

- Multi-modal learning signals in a bidirectionally-connected, hierarchically-organized network

1. Spike-based Learning

- Bridging the gap between biology and cognition

Learning to Look: Integrating over a “dumb” signal





Feedback error is in ***sensory*** coordinates
but applied directly to ***motor*** signal



(dumb = wrong coords, but only sensory is available; Kawato et al)

Looking to Reach: Same Error Signal, Different Muscles



1. Target On



2. Fixate Target



3. Initial Reach



How Emer's Brain Works

1. Learning to Look

- How a “dumb” signal can be smart (Kawato et al)


1. Recognizing Objects

- Multi-modal learning signals in a bidirectionally-connected, hierarchically-organized network

1. Spike-based Learning

- Bridging the gap between biology and cognition

Invariant Object Recognition



- Hierarchy of increasing:
 - Feature complexity
 - Spatial invariance
- Strong match to RF's in corresponding brain areas

(Fukushima, 1980; Poggio, Riesenhuber, et al...)

3D Object Recognition Test




- From Google SketchUp Warehouse
- 100 categories
- 8+ objects per category
- 2 objects left out for testing
- $\pm 20^\circ$ horizontal depth rotation + 180° flip
- $0-30^\circ$ vertical depth rotation
- 14° 2D planar rotations
- 25% scaling
- 30% planar translations



- Objects in Car Category


- Objects in Fish Category




- Depth & lighting variations for one object




Generalization Results: 92.8%



Bidirectional Multimodal Training




Benefit of Bidirectionality: Noise Robustness



Noise matched to image spatial frequency structure.

Benefit of Bidirectionality: Noise Robustness




Cluttered Backgrounds



- Performance degrades significantly
- Need figure-ground segregation – in V2

Embodied Figure Ground Learning



- Binocular foveation → disparity → 3D depth
- 3D depth trains 2D depth/occlusion cues

How Emer's Brain Works

1. Learning to Look

- How a “dumb” signal can be smart (Kawato et al)

1. Recognizing Objects

- Multi-modal learning signals in a bidirectionally-connected, hierarchically-organized network

1. Spike-based Learning

- Bridging the gap between biology and cognition

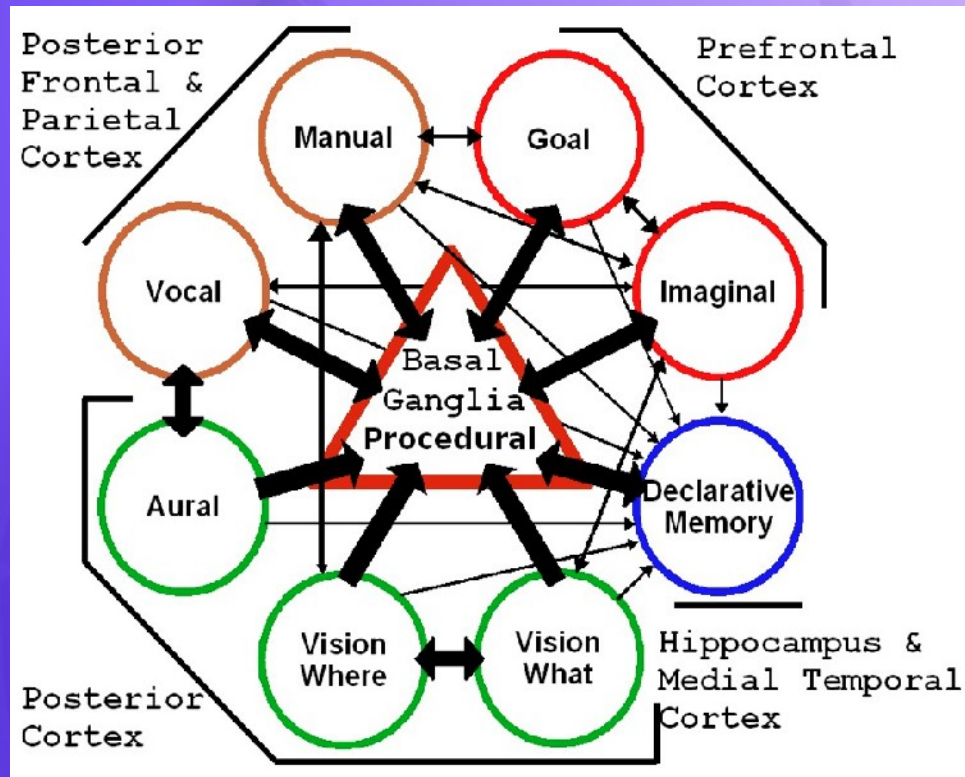
Embodiment = Rich Continuous Source of Learning Signals

- Every gaze is an automatic training signal
- Multiple modalities driven by same environment
 - Visual fixation trains motor reaching
 - Object foveation \leftrightarrow naming mutual training
- Binocular foveation \rightarrow disparity \rightarrow 3D depth \rightarrow trains 2D depth \rightarrow figure-ground segregation \rightarrow object recognition in cluttered scenes

The Common Sense Problem

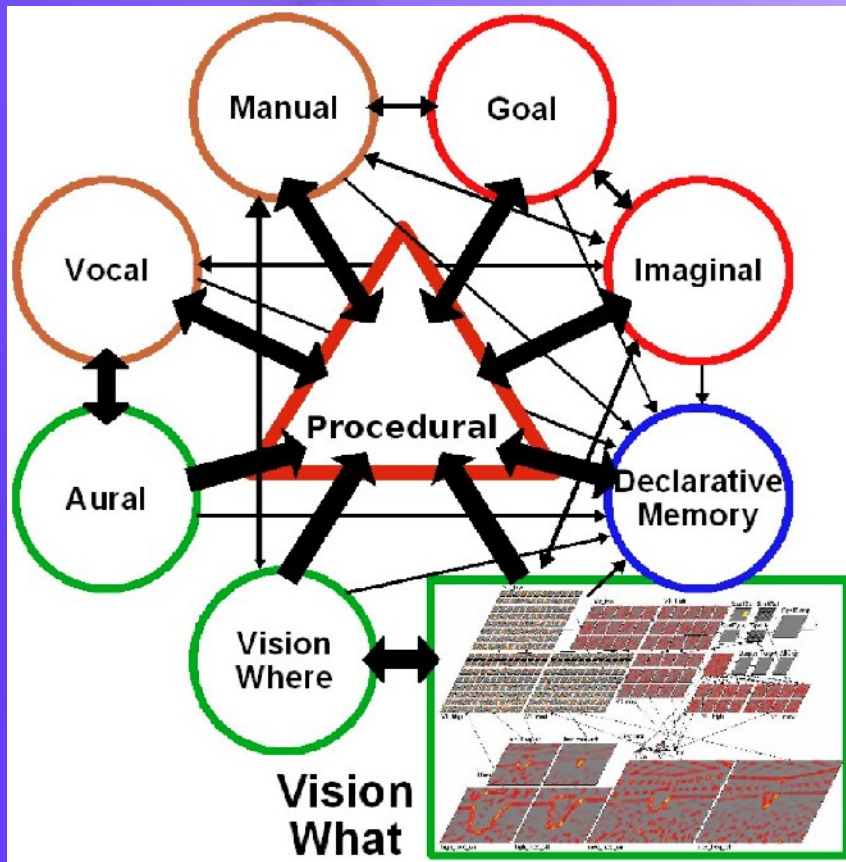
- Thesis: Human common sense depends on intensive learning about physical world
- More abstract cognitive functions build upon this sensory-motor foundation (Barsalou et al)
- Need this common sense to truly understand natural language and interact with people

SAL = Synthesis of ACT-R & Leabra




- Very close mapping between ACT-R and Leabra archs
- Productive ongoing collaboration to explore & develop SAL.

Simple Modular SAL



- Replace ACT-R module with Leabra visual system
- Achieves mutual benefits quickly
- But not as satisfying as a fuller synthesis..

The Leabra Algorithm



Same algorithm (and mostly params) can simulate ~100 different cognitive phenomena!